

# **Explainable Gridworld Navigation Using ILP**

Ayush Salunke

Ayushi Arora

Mentor: Youssef Mahmoud Youssef

# Problem Statement

This project focuses on building an explainable Gridworld navigation agent.

- Transparent decision-making is essential in AI systems
- Many learning-based models operate as black boxes
- Gridworld navigation is a simple but representative planning problem

## **Objective:**

- Learn how an agent should choose actions
- Avoid manually defining navigation rules
- Ensure all decisions remain explainable

# Popper for ILP

- Popper is an inductive logic programming system. Popper combines logical reasoning with machine learning to induce rules from examples and background knowledge. [\[1\]](#)
- Instead of manually defining navigation rules, the system employs the Popper ILP framework to induce a symbolic action-selection rule from examples and background knowledge.
- The learned rule is subsequently integrated into a Prolog-based planner to generate paths from a start position to a goal position while avoiding obstacles.
- Popper requires three input files:
  - an examples file
  - a background knowledge (BK) file
  - a bias file

# Gridworld environment

- $9 \times 9$  grid
- Start position: (0,0)
- Goal position: (3,8)
- Obstacles placed manually
- Allowed actions: up, down, left, right

8		red			green		red	red	red
7		red			blue		blue	blue	blue
6					red		white	white	red
5					white		red	red	white
4					red	red	blue	blue	red
3	red	red			blue	blue	blue	red	blue
2			blue	blue	white	white	white	red	white
1	red		blue	red			red		
0	green	blue	blue	white	red	red	white		
	0	1	2	3	4	5	6	7	8

# Positive & Negative examples

- Positive examples define correct behavior
- Negative examples define forbidden actions
- An examples file contains positive and negative examples of the relation you want to learn
- For example:

```
?- best_action((8,5), up).  
true.  
  
?- best_action((6,4), right).  
true.  
  
?- best_action((4,2), right).  
true.
```

Positive Example

```
?- best_action((6,3), up).  
false.  
  
?- best_action((9,5), up).  
false.  
  
?- best_action((0,0), left).  
false.
```

Negative Example

# Learning Framework

- Learning performed using Inductive Logic Programming (ILP)
- Implemented with the Popper framework
- Learns symbolic rules instead of numeric parameters
- Target predicate:

best\_action(Position, Action)

- Output is a logical Prolog clause
  - Which gives a link between the best action to take and the actions available.

```
***** SOLUTION *****
Precision:1.00 Recall:0.60 TP:6 FN:4 TN:10 FP:0 Size:2
best_action(V0,V1):- improving_move(V0,V1).
*****
```

# Explainable Decision Making

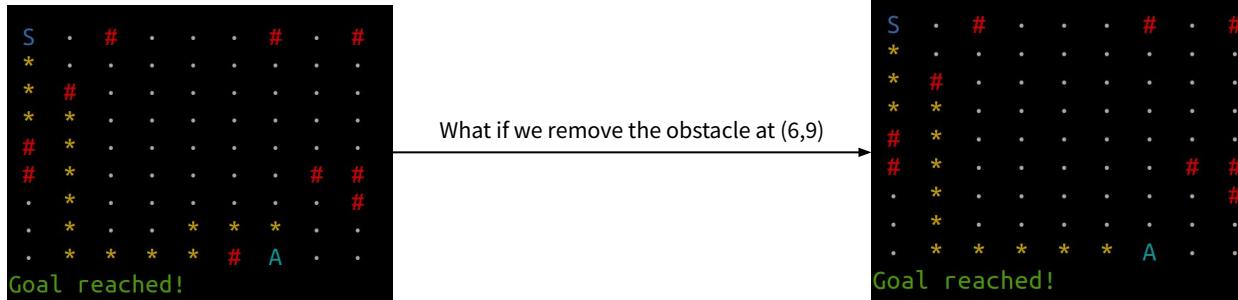
- Decisions represented as logical rules
- Learned model is symbolic and human-readable
- No neural networks or numeric parameters
- Action selection based on distance reduction
- Each action can be logically justified
- Explanation example:

“This action reduces the Manhattan distance to the goal.”

# Path from Start to Goal

# Counterfactual Reasoning

- This is a simple “What-if?” analysis.
- It helps enable us a causal reasoning like, “If a certain fact is changed in the world, how would it affect the output?”
- Supported counterfactuals for our program are:
  - What if we remove an obstacle
  - What if we add an obstacle
  - What if we change the goal
  - What if we change the start position



\*Note: for now we have just implemented this in the random gridworld program only for testing purposes

# Limitations & Future Work

## Limitations:-

- Currently there are no global planning algorithm.
- The current greedy heuristic (manhattan distance) is not an optimal algorithm.
- Popper learning does not work with a randomly generating gridworld.

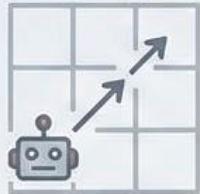
## Future work:-

- We can work on a different search algorithm like A\*.
- We can allow diagonal movements for the agent.
- We can also implement a bigger gridsearch or provide a real time UI to track the agents movements.
- Integrate popper ilp with a random start, goal positions along with random obstacles.

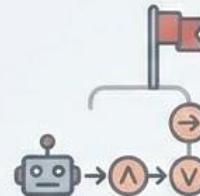
# **Thank You!**

Any Questions?

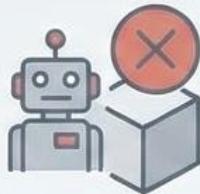
# Problem Statement



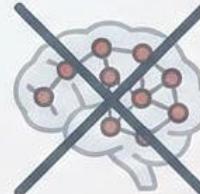
**Goal:** Navigate an agent in a grid world



Reach goal using  
explainable logic rules



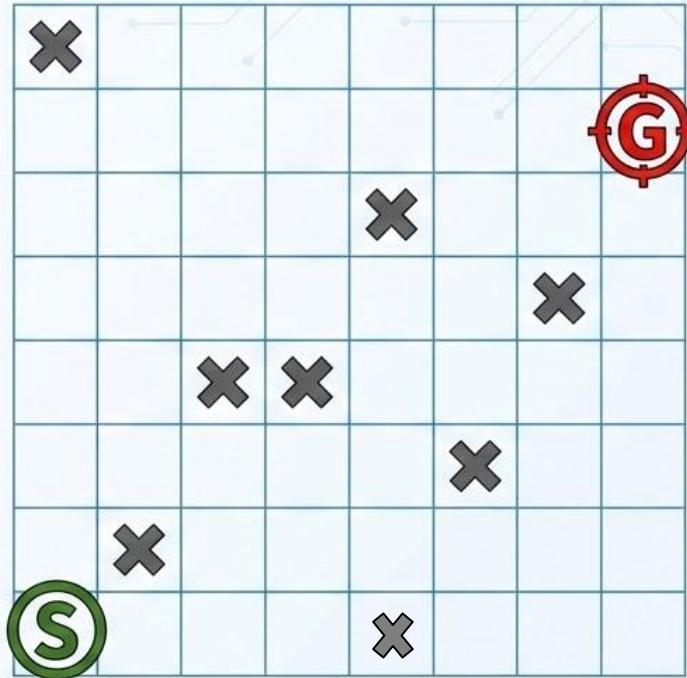
Avoid obstacles



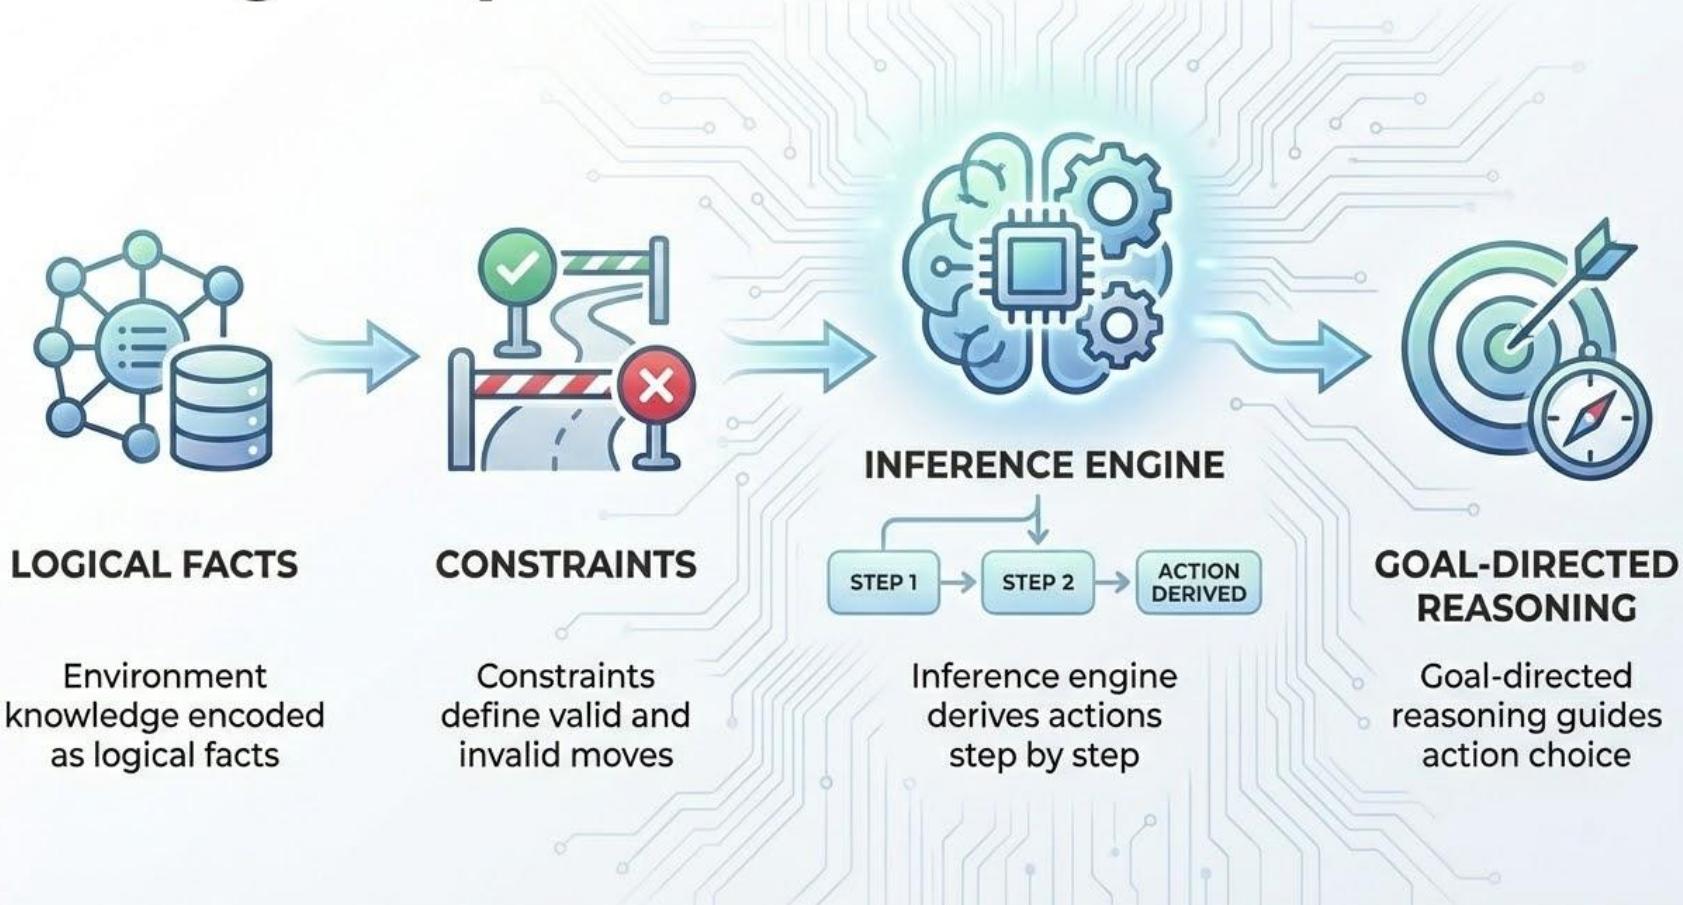
No machine learning,  
no training

# Gridworld environment

- 9 × 9 grid
- Start position: (0,0)
- Goal position: (8,7)
- Obstacles placed manually
- Allowed actions: up, down, left, right



# Logic Representation and Inference



# Explainable Decision Making



## EXAMPLE OUTPUT:

Chosen action: right

→ Reason: reduces distance to goal ←

# Verification: Positive & Negative examples



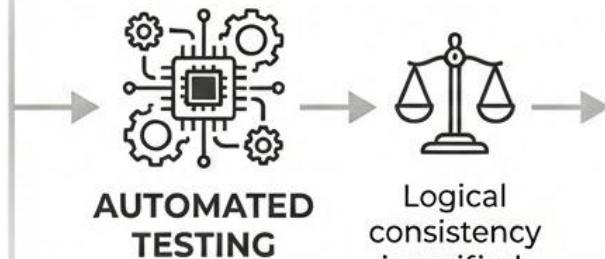
## POSITIVE EXAMPLES

Define correct behavior.  
Establish the desired path.



## NEGATIVE EXAMPLES

Define forbidden actions.  
Prevent invalid moves.



## PROLOG CODE EXAMPLE

```
?- best_move((1,0), right).  
\u001b[1;32mtrue.\u001b[0m
```

```
?- best_move((1,0), up).  
\u01b[1;31mfalse.\u001b[0m
```

# Motivation & Problem Statement

In many AI systems, decisions are difficult to interpret. Explainability requires that an agent's behavior be transparent and traceable and Prolog provides a natural framework for explainable decision-making because all reasoning is symbolically defined and fully visible.

This project builds an explainable Gridworld agent.

Create a small Gridworld environment in Prolog where an agent must reach a goal while avoiding obstacles. Instead of learning through trial-and-error, the agent uses logical rules to choose actions.

# Gridworld environment

- $9 \times 9$  grid
- Start position: (0,0)
- Goal position: (8,7)
- Obstacles placed manually
- Allowed actions: up, down, left, right

